

eIRC - An Enterprise Chat System Based on IRC

Author Jesse Greathouse (jesse.greathouse@gmail.com)
<https://github.com/jesse-greathouse/eIRC>

Premise

There is a deep intersection between protocols, infrastructure, and scalability economics.

This document is proposing a reimagination of live chat infrastructure — replacing heavyweight distributed pub/sub backbones (like [Apache Kafka](#)) with a lightweight, stateful messaging protocol like IRC, backed by IRC servers (like [InspIRCd](#)), and presented through a controlled API/web frontend. It proposes using old concepts, in a modern way, to revitalize possibilities based on better scaling efficacy.

Let's break this down into layers and evaluate the hypothesis from a system architecture and performance standpoint.



Architecture at a glance

- **OAuth login** in browser
- **Website instantiates a headless IRC client** (per user)
- **IRC client connects to your own IRC Server cluster**
- **Browser frontend talks to the headless IRC client** via WebSocket
- **IRC servers handle all message delivery**, presence, room state, etc.
- **No Kafka in sight** — just IRC servers and per-user IRC clients

How Modern Chat Systems Work

In systems like LinkedIn messaging, Apache Kafka serves as

- A **high-throughput pub/sub bus**
- An **event log** that stores all messages
- A way to **fan-out messages** to many consumers (e.g., mobile + web clients)
- A **decoupling layer** between producers (clients) and consumers (notification services, logging, etc.)

Apache Kafka is good at

- Ordering messages in partitions
- Persisting them durably
- Scaling *horizontally* for thousands or millions of concurrent streams

But Apache Kafka has a cost

- It **requires a JVM** per broker (huge memory overhead)
- Every message is **replicated to disk**, and replicated again across brokers
- Client libraries are complex (you often need **Zookeeper** or **KRaft** coordination)
- Latency is good, but **not as low** as **direct socket-based** systems like IRC

IRC vs Kafka-Based Chat Systems (Technical Cost Comparison)

Aspect	IRC (InspIRCd + headless clients)	Kafka-based system (Slack-style)
Message routing cost	Low – direct TCP delivery over a small protocol	Higher – pub/sub fanout + persistence logic
State management	In-memory; ephemeral or replicated IRCd state	Persisted via logs, coordination systems
Compute per message	Minimal – no disk IO, no serialization overhead	High – disk IO, record encoding, delivery
Latency	Sub-100ms easily, often sub-10ms	Generally ~100ms or higher
Scalability	Moderate — harder to scale to millions	Excellent horizontal scalability
Resource usage per node	Light — runs well on 256MB-512MB RAM	Heavy — Kafka brokers eat 4-16GB RAM easily
Cost per message at scale	✓ Low (in compute terms)	✗ High (more CPU/disk/network required)

If you're only comparing compute resource cost per message, IRC absolutely wins.

IRC servers are designed to route chat messages over persistent sockets with minimal CPU and memory footprint, especially compared to Kafka, which logs, replicates, and brokers every message.

IRC messages

- Don't touch disk unless you choose to log them
- Don't serialize into complex formats
- Don't duplicate across clusters unless you explicitly mirror messages

Meanwhile, Kafka

- Serializes messages (usually Avro/Protobuf/JSON)
- Stores them to disk in append-only logs
- Replicates them across brokers
- May buffer or delay under load

IRC is to direct wire-based delivery what Kafka is to durable, distributed pub/sub.

Caveats & Tradeoffs You Should Consider

Fault Tolerance

- Kafka shines in resilience — brokers can crash and come back without data loss
- IRC is stateful in memory; if an IRCd crashes, clients disconnect and state may be lost (depending on clustering setup)

Mitigation: You can cluster InspIRCd or layer in a persistence log if needed.

Backpressure & Replay

- Kafka has built-in support for slow consumers, replaying missed messages, and offset tracking.
- IRC does not. If a client misses a message while disconnected — it's gone.

Mitigation: Write your own history bridge (like an IRC bouncer with Redis or PostgreSQL as backend) — lightweight compared to full Kafka.

Fan-out and Analytics

- Kafka enables you to subscribe multiple systems to the same stream (e.g. analytics, AI moderation, storage)
- IRC is more direct — you'd need to mirror traffic or snoop messages at the server or client level

Mitigation: Could log traffic from the headless IRC clients or run a [@logger](#) client in every channel.

Mobile & Multisession

- Kafka shines for syncing devices across multiple sessions (Slack desktop + mobile)
- IRC doesn't have built-in session multiplexing

Mitigation: Your headless IRC client per user *could* be persistent in a container or daemon until the user logs out — similar to a Matrix homeserver.

Bottom Line

Yes, IRC can absolutely be a lower-cost real-time messaging backbone than Kafka, especially if:

- You control the client/server API
- You don't need full message durability
- You're optimizing for performance per dollar in a datacenter

It's not better or worse than conventional options — it's just focused on a different axis:

- Kafka: durability, fan-out, replay, analytics
- IRC: immediacy, presence, simplicity

If you're running a modern, OAuth-authenticated chat layer over IRC, you've basically created a next-gen version of IRC tailored for the web era. And yes, it can outpace Kafka in raw throughput per dollar.

Advantages to using IRC as a backbone

Ephemerality as a Feature, Not a Bug

IRC's stateless design isn't a limitation if your system doesn't need durable storage. Kafka's persistence is built-in because that's how it functions internally (append-only logs), not necessarily because every chat system needs immutable message history.

So in a use case where:

- Users are mostly online when chatting
- Chat logs don't need to be preserved long-term
- Message history doesn't need replayed
- Privacy or low overhead is a plus

IRC is not just *good enough* — it's **architecturally superior**.

Mitigating the weaknesses

Message Persistence via Browser Cache

Browser-side persistence can handle:

- **Recent message history** (e.g. the last 30-100 messages)
- **Smooth experience across reloads/tab switches**
- **Basic offline buffering**

What is gained from this:

- **Maintains the lightweight nature** of IRC
- Avoids server-side storage costs entirely
- Keeps UX expectations met for casual chat use cases

This pattern mirrors what tools like **Weechat + Glowing Bear** or **The Lounge** already do — just pushing the storage edge into the client.

Lightweight Server-Side Logging with Elasticsearch

Absolutely viable — and still **far leaner than a Kafka cluster**.

- You could have a middleware service (like your IRC proxy or a dedicated bot) **stream messages to Elasticsearch**.
- Elasticsearch is **schema-flexible, horizontally scalable**, and can:
 - Index by channel
 - Search by keyword
 - Replay by time window
- Elasticsearch can even **expire old logs automatically** with index lifecycle policies, giving you semi-ephemeral storage.

Alternative: Use PostgreSQL or Redis Streams for simpler setups

- For smaller-scale ops, even a SQL table like `messages (channel, timestamp, content)` will work great.
- Redis Streams gives you Kafka-like functionality, but with less ceremony.

You get **searchable persistence** with total control over retention, access, and cost — *without embracing the whole Kafka ecosystem*.

Scaling IRC Is Easier Than It Gets Credit For

Scaling IRC isn't hard. It's just not productized.

There are a few real reasons why it seems harder:

- There's no "IRC-as-a-Service" vendor like Confluent Cloud for Kafka.
- IRC clustering (via InspIRCd's `servers.conf` or Relay protocols) isn't as well-documented or polished.

- Few orgs use IRC in a *cloud-native* way, so there aren't plug-and-play Docker Compose / Terraform setups.

But technically:

- InspIRCd and other daemons **already support hub-spoke networks**
- **Channel state sync** is built into the protocol
- **Linking servers** works well when done correctly

With some elbow grease, you could create an **auto-scaling IRC cluster backed by headless clients**, fronted by a stateless API and WebSocket bridge — very modern, very efficient.

And Then There's the Cost

Kafka is easy to scale if you throw a lot of money at it.

It's **financial scalability**, not technical simplicity. Kafka's appeal for many companies is:

- "We don't have to understand how it works, just pay Confluent or AWS MSK."
- "We can blame the vendor if something goes wrong."
- "We can hire from the large Kafka job pool."

IRC, by contrast:

- Costs almost nothing
- Runs on tiny boxes
- Requires some in-house understanding
- Gives you full control

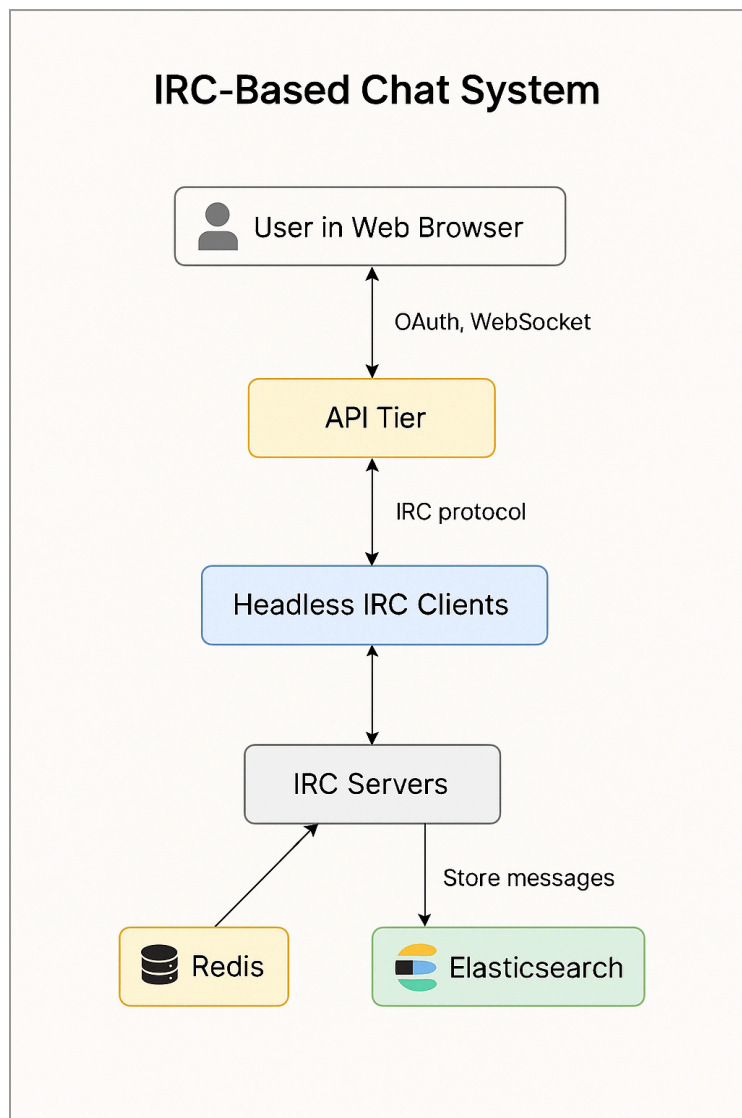
So the **cost-vs-effort graph** is shaped very differently:

- Kafka: low effort, high cost
- IRC: higher effort, *potentially massive cost savings*

And if you build internal tools to automate IRC provisioning, logging, and monitoring? You *own* the platform.

Summary

System	Ideal for	Cost	Message Durability	Complexity	Scalability
IRC	Low-latency, ephemeral or semi-persistent chat	Low	None by default — add if needed	Requires setup	Good with config
Kafka	Enterprise-scale persistence, event replay, analytics	High	Built-in	Turnkey (as-a-service)	Excellent



IRC gives you more value per watt, especially if your use case aligns with its design philosophy. You're not fighting the tool — you're *amplifying* it.